

# Рубежный контроль №2

Садыков Марк, ИУ5-63Б

**Тема: Технологии использования и оценки моделей машинного обучения.**

In [4]:

```
from sklearn.linear_model import LogisticRegression
from sklearn.svm import LinearSVC
from sklearn.naive_bayes import MultinomialNB, ComplementNB, BernoulliNB, CategoricalNB
from sklearn.metrics import accuracy_score
from sklearn.datasets import fetch_20newsgroups
from sklearn.feature_extraction.text import TfidfVectorizer
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(style="ticks")
from IPython.display import set_matplotlib_formats
set_matplotlib_formats("retina")
pd.set_option("display.width", 70)
```

In [5]:

```
data_train = fetch_20newsgroups(subset='train', remove=('headers', 'footers'))
data_test = fetch_20newsgroups(subset='test', remove=('headers', 'footers'))
```

In [7]:

```
data_train.target.shape
```

Out[7]:

```
(11314,)
```

In [9]:

```
data.dtypes
```

Out[9]:

Employee_Name	object
EmpID	float64
MarriedID	float64
MaritalStatusID	float64
GenderID	float64
EmpStatusID	float64
DeptID	float64
PerfScoreID	float64
FromDiversityJobFairID	float64
PayRate	float64
Termd	float64
PositionID	float64
Position	object
State	object
Zip	float64
DOB	object
Sex	object
MaritalDesc	object
CitizenDesc	object
HispanicLatino	object
RaceDesc	object
DateofHire	object
DateofTermination	object
TermReason	object
EmploymentStatus	object
Department	object
ManagerName	object
ManagerID	float64
RecruitmentSource	object
PerformanceScore	object
EngagementSurvey	float64
EmpSatisfaction	float64
SpecialProjectsCount	float64
LastPerformanceReview_Date	object
DaysLateLast30	float64
dtype:	object

In [8]:

```
vectorizer = TfidfVectorizer()  
vectorizer.fit(data_train.data + data_test.data)
```

Out[8]:

```
TfidfVectorizer(analyzer='word', binary=False, decode_error='strict',  
                dtype=<class 'numpy.float64'>, encoding='utf-8',  
                input='content', lowercase=True, max_df=1.0, max_features=None,  
                min_df=1, ngram_range=(1, 1), norm='l2', preprocessor=None,  
                r=None, smooth_idf=True, stop_words=None, strip_accents=None,  
                e, sublinear_tf=False, token_pattern='(?u)\\b\\w\\w+\\b',  
                tokenizer=None, use_idf=True, vocabulary=None)
```

In [9]:

```
X_train = vectorizer.transform(data_train.data)  
X_test = vectorizer.transform(data_test.data)  
  
y_train = data_train.target  
y_test = data_test.target
```

In [11]:

```
def test(model):  
    print(model)  
    model.fit(X_train, y_train)  
    ac = accuracy_score(y_test, model.predict(X_test))  
    print("accuracy:", ac)  
    return ac
```

In [12]:

```
test(LogisticRegression(solver='lbfgs', multi_class='auto'))
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,  
                   intercept_scaling=1, l1_ratio=None, max_iter=100,  
                   multi_class='auto', n_jobs=None, penalty='l2',  
                   random_state=None, solver='lbfgs', tol=0.0001, verbose=0,  
                   warm_start=False)  
accuracy: 0.774429102496017
```

Out[12]:

0.774429102496017

In [13]:

```
test(LinearSVC())
```

```
LinearSVC(C=1.0, class_weight=None, dual=True, fit_intercept=True,  
          intercept_scaling=1, loss='squared_hinge', max_iter=1000,  
          multi_class='ovr', penalty='l2', random_state=None, tol=0.  
0001,  
          verbose=0)  
accuracy: 0.8048327137546468
```

Out[13]:

0.8048327137546468

In [14]:

```
test(MultinomialNB())
```

```
MultinomialNB(alpha=1.0, class_prior=None, fit_prior=True)  
accuracy: 0.72623473181094
```

Out[14]:

0.72623473181094

In [15]:

```
test(ComplementNB(alpha=0.3))
```

```
ComplementNB(alpha=0.3, class_prior=None, fit_prior=True, norm=False)  
accuracy: 0.812931492299522
```

Out[15]:

0.812931492299522

In [16]:

```
test(BernoulliNB())
```

```
BernoulliNB(alpha=1.0, binarize=0.0, class_prior=None, fit_prior=True)  
accuracy: 0.5371747211895911
```

Out[16]:

0.5371747211895911

In [17]:

```

balpa=0
bac=0
i=1
while i>0:
    res=test(ComplementNB(alpha=i))
    # print(res)
    if res>bac:
        balpa=i
        bac=res
    i-=0.1
print(balpa, " ", bac, "\n")

```

```

ComplementNB(alpha=1, class_prior=None, fit_prior=True, norm=False)
accuracy: 0.8089484864577802
ComplementNB(alpha=0.9, class_prior=None, fit_prior=True, norm=False)
accuracy: 0.8102761550716941
ComplementNB(alpha=0.8, class_prior=None, fit_prior=True, norm=False)
accuracy: 0.8101433882103027
ComplementNB(alpha=0.7000000000000001, class_prior=None, fit_prior=True,
              norm=False)
accuracy: 0.8117365905469994
ComplementNB(alpha=0.6000000000000001, class_prior=None, fit_prior=True,
              norm=False)
accuracy: 0.8120021242697822
ComplementNB(alpha=0.5000000000000001, class_prior=None, fit_prior=True,
              norm=False)
accuracy: 0.8117365905469994
ComplementNB(alpha=0.40000000000000013, class_prior=None, fit_prior=True,
              norm=False)
accuracy: 0.8127987254381307
ComplementNB(alpha=0.30000000000000016, class_prior=None, fit_prior=True,
              norm=False)
accuracy: 0.812931492299522
ComplementNB(alpha=0.20000000000000015, class_prior=None, fit_prior=True,
              norm=False)
accuracy: 0.811603823685608
ComplementNB(alpha=0.10000000000000014, class_prior=None, fit_prior=True,
              norm=False)
accuracy: 0.8101433882103027
ComplementNB(alpha=1.3877787807814457e-16, class_prior=None, fit_prior=True,
              norm=False)
accuracy: 0.6955655868295274
0.30000000000000016    0.812931492299522

```

```

/home/mark/.local/lib/python3.7/site-packages/sklearn/naive_bayes.py:507: UserWarning: alpha too small will result in numeric errors, setting alpha = 1.0e-10
  'setting alpha = %.1e' % _ALPHA_MIN)

```

In [18]:

```

balpa=0
bac=0
i=0.2
while i>0:
    res=test(MultinomialNB(alpha=i))
    # print(res)
    if res>bac:
        balpa=i
        bac=res
    i-=0.02
print(balpa," ",bac,"\n")

```

```

MultinomialNB(alpha=0.2, class_prior=None, fit_prior=True)
accuracy: 0.7738980350504514
MultinomialNB(alpha=0.18000000000000002, class_prior=None, fit_prior
=True)
accuracy: 0.7769516728624535
MultinomialNB(alpha=0.16000000000000003, class_prior=None, fit_prior
=True)
accuracy: 0.7804036112586299
MultinomialNB(alpha=0.14000000000000004, class_prior=None, fit_prior
=True)
accuracy: 0.782793414763675
MultinomialNB(alpha=0.12000000000000004, class_prior=None, fit_prior
=True)
accuracy: 0.7851832182687202
MultinomialNB(alpha=0.10000000000000003, class_prior=None, fit_prior
=True)
accuracy: 0.7886351566648965
MultinomialNB(alpha=0.08000000000000003, class_prior=None, fit_prior
=True)
accuracy: 0.7908921933085502
MultinomialNB(alpha=0.060000000000000026, class_prior=None, fit_prio
r=True)
accuracy: 0.7943441317047265
MultinomialNB(alpha=0.040000000000000002, class_prior=None, fit_prior
=True)
accuracy: 0.7987254381306426
MultinomialNB(alpha=0.020000000000000002, class_prior=None, fit_prior
=True)
accuracy: 0.8023101433882103
MultinomialNB(alpha=2.0816681711721685e-17, class_prior=None, fit_pr
ior=True)
accuracy: 0.7458842272968667
0.020000000000000002    0.8023101433882103

```

```

/home/mark/.local/lib/python3.7/site-packages/sklearn/naive_bayes.p
y:507: UserWarning: alpha too small will result in numeric errors, s
etting alpha = 1.0e-10
'setting alpha = %.1e' % _ALPHA_MIN)

```

In [19]:

```
balpa=0
bac=0
i=0.1
while i>0:
    res=test(BernoulliNB(alpha=i))
    #     print(res)
    if res>bac:
        balpa=i
        bac=res
    i-=0.01
print(balpa, " ", bac, "\n")
```

```

BernoulliNB(alpha=0.1, binarize=0.0, class_prior=None, fit_prior=True)
accuracy: 0.6435209771640998
BernoulliNB(alpha=0.090000000000000001, binarize=0.0, class_prior=None,
              fit_prior=True)
accuracy: 0.6456452469463622
BernoulliNB(alpha=0.080000000000000002, binarize=0.0, class_prior=None,
              fit_prior=True)
accuracy: 0.6473712161444504
BernoulliNB(alpha=0.070000000000000002, binarize=0.0, class_prior=None,
              fit_prior=True)
accuracy: 0.650292087095061
BernoulliNB(alpha=0.060000000000000002, binarize=0.0, class_prior=None,
              fit_prior=True)
accuracy: 0.6524163568773235
BernoulliNB(alpha=0.050000000000000002, binarize=0.0, class_prior=None,
              fit_prior=True)
accuracy: 0.6548061603823686
BernoulliNB(alpha=0.0400000000000000015, binarize=0.0, class_prior=None,
              fit_prior=True)
accuracy: 0.6589219330855018
BernoulliNB(alpha=0.0300000000000000013, binarize=0.0, class_prior=None,
              fit_prior=True)
accuracy: 0.6631704726500266
BernoulliNB(alpha=0.020000000000000001, binarize=0.0, class_prior=None,
              fit_prior=True)
accuracy: 0.6684811471056824
BernoulliNB(alpha=0.010000000000000001, binarize=0.0, class_prior=None,
              fit_prior=True)
accuracy: 0.6743228890069038
BernoulliNB(alpha=1.0408340855860843e-17, binarize=0.0, class_prior=None,
              fit_prior=True)
accuracy: 0.7132235793945831
1.0408340855860843e-17    0.7132235793945831

```

```

/home/mark/.local/lib/python3.7/site-packages/sklearn/naive_bayes.py:507: UserWarning: alpha too small will result in numeric errors, setting alpha = 1.0e-10
  'setting alpha = %.1e' % _ALPHA_MIN)

```

In [ ]: